

CS 61A Final  
Spring 2010

Question 1:

What is the order of growth in time of the following function? Check all that apply.

```
(define (foo x y)
  (cond ((= x 0) y)
        ((= x 1) (foo (- x 1) (* x y)))
        (else (foo (- x 1) (foo 1 (+ x y))))))
```

At first blush, this might seem exceptionally difficult because of the

recursive call as one of the arguments to FOO in the ELSE clause. However,

notice that the logic of FOO only depends on the value of X; when X is 0,

return Y, when X is 1, do something else, etc. The recursive call to FOO:

(foo 1 (+ x y)) in the ELSE clause has X set to 1 for the next call, which

causes it to call (foo (- 1 1) (\* 1 y)) for the next call, which will

eventually return Y. That means that the recursive call (foo 1 (+ x y))

always takes constant time, so we can essentially ignore it. Then we're

left with X decreasing by 1 at each recursive call, so it takes  $O(x)$  time.

Scoring: 2pts, all or nothing

Question 2:

When you reserve a book at the library, you can give them your e-mail address, so they can send you an e-mail when your book comes

in. This is an example of:

a callback

The answer would have been a bit too obvious if we said the library "calls you back", so we used e-mail instead. The callback here is your e-mail address; once something happens that will require them to "execute" the callback (your book comes back in), the original thing that requested the information (you) gets notified.

A particular kindergarten class has the rule that only the person holding a certain special sock can talk. The sock is an example of:

a mutex

Everyone is waiting on the sock to be able to talk, so it acts as a mutex in this case. Note that it couldn't be a deadlock; the person speaking could simply hand off his sock if someone else ever had to talk!

Scoring: 2pts each, all or nothing

Question 3:

(Based on a true story!) A certain model of fighter plane had an abnormally high number of crashes on landing -- blamed on pilot error. It turned out the controls for the flaps and the landing gear were right next to each other, causing pilots to accidentally retract the landing gear when they meant to extend the flaps to slow the plane down. The fix? A

round knob on the landing gear lever, and a rectangular one on the flaps control. The change stopped these accidents almost completely.

How is this story similar to the case of the Therac-25?

Answer: "Operator error" was a result of a poor user interface.

While it's true that it is the pilot that was at fault, the core misunderstanding that caused the problem is at the user interface. The two controls are operated at nearly the same time, and have the same look; naturally problems would occur! The Therac-25 also had user interface problems, namely cryptic error messages that the operator would simply skip once he/she got used to it, and fields that allowed you to type over previous settings without updating them.

The other answers:

The problem was easy to reproduce.

True of the airplane, not true of the Therac-25.

Incorrect concurrency was one of the main issues.

True of the Therac-25, not true of the airplane.

The makers of the machines actively tried to hide the problem.

True of neither.

Scoring: 2pts, all or nothing

Question 4:

Given the following code:

```
(define x 10)
(define y 10)

(define foo
  (let ((x 5))
    (lambda (y) (+ x y))))

(foo 7)
```

What is the result of the final expression using lexical scope? 12

What is the result of the final expression using dynamic scope? 17

Which does Scheme use? lexical scope

Note that FOO is defined as the return value of its LET; the LET here returns the LAMBDA whose lexical environment (right bubble) points to the frame created by the LET. In other words, in lexical scope, the X that the LAMBDA will access is the LET's X. So when you call (foo 7), X is 5, and Y is 7, giving you 12.

In dynamic scope, when we call FOO with 7, the X the body of the LAMBDA gets is the one in the current (global frame). X is 10, Y is 7, giving us 17.

Finally, the question about Scheme should've been a freebie if you've been paying attention in class :)

Scoring: 5pts - 2 for each blank and 1 for "Which does Scheme use?"

(If you managed to get the entire thing reversed, we still gave you a single point for recognizing which answer Scheme should get.)

Question 5:

Here is the code with the blanks filled in:

```
; Assume x, y, and z have already been defined.
> (define a (append (list x y) z))
a
> (define three (list 3))
three
> a
((1 2) (1 2) 1 2)

> (set-cdr! (cdr z) THREE)
okay
> a
((1 2) (1 2) 1 2 3)

; either (set-cdr! X (CDR Z)), or
> (set-cdr! (CDR X) THREE)
okay
> a
((1 2 3) (1 2 3) 1 2 3)

> (set-car! THREE 6)
okay
> (set-car! X 4)
okay
> a
((4 2 6) (4 2 6) 1 2 6)
```

The trick was to realize that while X, Y, and Z all print as (1 2), X and Y are EQ?, meaning they point to the same box in a box-and-pointer diagram.

After that it was a relatively simple matter of keeping your cars and cdrs

straight.

Most of the blanks had more than one correct answer -- for example, you could have said `(cdar a)` instead of `(cdr x)`. Except for the "either-or" comment above, though, there was only one correct /value/ (usually a pair) that could go in each slot. We would have given credit even to answers like `(set-cdr! (cdadr a) (cddddr a))`, which is just `(set-cdr! (cdr y) three)`.

Scoring: 1pt per blank. Because some blanks were harder than others, this kept people from losing all points on this problem.

#### Question 6:

To answer the question, you were not supposed to solve the puzzle, but to translate the puzzle into an evaluation through `ambeval`.

In the puzzle, you were supposed to find the order in which Alice, Bob, Carol and Dave arrived. This was encoded into the variables `a`, `b`, `c` and `d` which could take values from 1 to 4.

The code has three parts to it: binding the variables to the ambiguous values (1-4), checking all the conditions posed in the puzzle (5-10) and then return the solution (that part was already given to you) in proper format.

The first of the seven bullet points in the quiz is just general

information, the next 6 translate to a (require ...  
) expressions each:

```
(let
  ((a (amb 1 2 3 4)) ; 1
   (b (amb 1 2 3 4)) ; 2
   (c (amb 1 2 3 4)) ; 3
   (d (amb 1 2 3 4))) ; 4
  (require (distinct? (list a b c d))) ; 5
  (require (or (> c a) (> c b))) ; 6
  (require (not (and (< a c) (< b c)))) ; 7
  (require (> d a)) ; 8
  (require (not (= d 4))) ; 9
  (require (= b 1)) ; 10
  (list (list 'Alice a) (list 'Bob b)
        (list 'Carol c) (list 'Dave d)))
```

Alternatively, you can drop line 9 by changing 4 to  
(d (amb 1 2 3)) and  
you can drop line 10 by changing 2 to (b (amb 1)) o  
r just (b  
1). Instead of (require ...), you can write (if (no  
t ...) (amb)).

Scoring: 5 points in total.

Lines 1-4 were worth 2 points in total. Each line 5  
-10 was worth half  
a point each. The result was rounded down.

Remark: The puzzle actually has no a solution!. So  
just writing  
((a (amb)) (b (amb)) (c (amb)) (d (amb))) gives the  
right result, but  
it would not be worth credit.

Question 7:

Here is the mystery code again:

```
(define (helper s n)
```

```

    (if (= n 0)
        the-empty-stream
        (cons-stream (stream-car s)
                      (helper (stream-cdr s) (- n 1)
                              )
                      )
    ))))

(define (mystery s)
  (define (foo n)
    ; like STREAM-APPEND, but takes a stream promise instead
    (stream-append-delayed (helper s n)
                           (delay (foo (+ n 1)))))
  )
  (foo 1))

```

Taking the hint, let's see what HELPER does. If N is 0, it returns the empty stream. Otherwise, it takes the first element of the stream, CONS-STREAMed with HELPER of the rest of the stream. The recursive call to HELPER decreases N by 1, so it will take N calls to reach the base case. So HELPER just takes the first N elements of the stream S. Notice that THE-EMPTY-STREAM is not an /element/ of the stream; it's the value that marks the end of a finite stream...just like the empty list marks the end of a list.

Now that we know what HELPER does, we can look at MYSTERY...which just calls FOO. FOO is going to append two streams: the first is (helper s n), and the second is (foo (+ n 1)). (Ignore the DELAY for right now.) So we're appending the first N elements of the input stream with (foo (+ n 1)). What does that stream start with? The first n+1 elements of the input stream. And so on.

So, MYSTERY takes a stream, and returns a stream of



the first element, the first two elements, the first three elements...appended together. The first ten elements are as follows:

```

1 1 4 1 4 9 1 4 9 16
^  ^^^  ^^^^^  ^^^^^^^^^
|      |           |               |
|      |           |               +-- (helper s 4)
|      |           +-- (helper s 3)
|      +-- (helper s 2)
+-- (helper s 1)

```

This was enough to get the answer, but one question remains: why the DELAY? Because Scheme is applicative order, and STREAM-APPEND isn't a special form! If we used regular STREAM-APPEND and a non-delayed recursive call, we would end up with infinite recursion trying to find (foo (+ n 1)).

Scoring: 5 points total  
 \* 2 for having subsequences that increased in size  
 \* 1 for having the correct numbers in those subsequences  
 \* 2 for making a flat stream instead of a stream of streams

Question 8:

First off, the solution we were looking for:

```

(define (fib-vector n)
  (let ((result (make-vector n 1))) ; initialize a
    ll values to 1
    (define (iter x)
      (if (>= x n)

```

```

        result
        (begin (vector-set! result x (+ (vector-r
ef result (- x 1))
                                         (vector-r
ef result (- x 2)) ))
              (iter (+ x 1)) )))
    (iter 2) ))

```

Notice how we use lexical scope to minimize the number of parameters that get passed around! The helper only needs to take one argument, the index of the current Fibonacci number to compute, and has access to the result vector and the limit anyway.

The interesting part of the problem is handling the base cases, which is what prompted us to announce that it was okay to assume  $n \geq 2$ . (The version above happens to work for any non-negative  $n$ .) There were three ways to do this:

- Like the solution above, use MAKE-VECTOR's implicit initialization.
- VECTOR-SET! the first two elements to 1, then call the helper.
- Have special cases in the helper for the first two elements.

Some people misinterpreted the question as suggesting the use of a memoized version of FIB or FIB-VECTOR -- i.e. numbers should n't be calculated more than once, ever! These were mostly graded on a case-by-case basis, but still had the same efficiency and usage criteria (e.g. using 1-based indexing still cost a point, etc.).

Scoring: 8 points total

- \* 1 for creating a vector of the proper size.

- \* 2 for getting the first two elements right.
- \* 1 for properly calling your helper procedure, based on how you initialized the first two elements.
- \* 2 for properly calculating the rest of the elements.
- \* 1 for knowing when to end the recursion (no off-by-one errors).
- \* 1 for remembering to return the vector at the end.

  

- \* -1 for using 1-based indexing; the first element of a vector is element 0.
- \* -2 for non-efficient uses of vectors, such as using the hypothetical procedure VECTOR-APPEND.
- \* -4 for calculating a specific Fibonacci number more than once, since the instructions specifically said not to. This included implementations that just called FIB.

### Question 9:

Write rules for RELATED-BY which take two people and a list of relationships.

The rules should match if and only if the two people are related through the sequence of relationships in the list.

Hint: What two people satisfy RELATED-BY if the list is empty?

The hint was intended to suggest the base case: if it doesn't take any relationships to get from person A to person B, then A and B must be the same person!

```
(assert! (rule (related-by ?self ?self ())))
```

How about the rest of it? Like most query system rules dealing with lists,  
we want to split the list of relationships up into the "first"  
relationship and the "rest" of the list. So the rule header we want is:

```
(assert! (rule (related-by ?a ?b (?first-rel . ?rest)) ...))
```

What are the conditions on this rule? That is, what must be true for the  
rule to match properly? We know that ?a has to be related to someone  
through the first relation, AND that that person is then related to ?b  
through the rest of the relations.

```
(assert! (rule (related-by ?a ?b (?first-rel . ?rest))  
              (and (?first-rel ?a ?someone)  
                   (related-by ?someone ?b ?rest))  
            )))
```

Actually, that's it! The "trick" was remembering that the query system  
isn't doing anything smart -- it's just pattern-matching. That means you  
can have two things related by ?first-rel, and ?first-rel will match  
PARENT or GRANDPARENT, or any other pattern in the query system.

The other important step was realizing that you needed to introduce a  
"middle" person, someone who was related to both ?a and ?b. This is  
also something we've seen before.

Some people took the additional step of adding (not (same ?a ?b)) as a  
condition. This isn't a bad idea, particularly

since SAME is a  
    "relationship" of sorts, and you could get your  
self into an infinite loop  
    that way: (related-by bart bart (same same same  
    same same ...))! But we  
    said not to worry about infinite loops.

We didn't take any points off for reversing the  
order of the relations,  
since we didn't specify it!

Scoring: 8pts total

- \* 2 for the base rule  
    (1 for an attempt)
- \* 6 for the recursive rule
  - 6 - correct
  - 5 - trivial errors
  - 4 - the idea: a recursive rule that goes one re  
lationship at a time and  
    introduces a "middle" person that's related  
to both ?a and ?b
  - 2 - an idea
  - 0 - other (hardcoding "parent" or "grandparent"  
in your rule(s), or trying  
    to use Scheme-style composition of function  
s)

Question 10:

Instead of the chain of definitions yielding to th  
e result of (best-song-ever 'artist) we could have  
written:

```
(
  (
    (
      generate-adt
      '(title artist)      ; let this be "names"
    )
    '(never gonna give you up) 'astley ; let this
    be "args"
```

```

    )
    'artist ; let this be "message"
)

```

We see that generate-adt's domain is lists (names, here (title artist)), its range is actually procedures (one of them is make-song), whose domain is a number of arguments (args, here (astley) and whose range is again procedures (one of them is best-song) whose domain is a word (message, here artist). It can be thought of as nested message-passing. It is three steps deep, so the solution should involve three lambdas. Furthermore, the body of the function has to figure out where the "message" appears in "names" with (position message names) and then return the corresponding element from args with (list-ref args (position message names)).

A more formal way to write the domain and range of this curried function is:  
generate-adt:

```

    lists ----> (objects^* ----> (strings ---> objects))
    names |---> ( args |-----> (message |--> (list
-ref args

    (position

    message

    names))))

```

And here is the scheme code:

```
(define generate-adt
  (lambda (names)
    (lambda args
      (lambda (message)
        (list-ref args (position message names))))))
```

or using the header given in the exam:

```
(define (generate-adt names)
  (lambda args
    (lambda (message)
      (list-ref args (position message names)
    ))))
```

Notice that in the above lambdas, there are no parentheses around args. This is intentional. Recall that scheme allows the following patterns for lambda:

lambda expression applied	what will variables bind to
<code>(lambda args args) 1 2 3)</code>	args = (1 2 3)
<code>(lambda (head . tail)) 1 2 3)</code>	head = 1, tail = (2 3)
<code>(lambda (a b c)) 1 2 3)</code>	a = 1, b = 2, c = 3

Hence the following is also a valid solution (except for the degenerate case of `(generate-adt '())`):

```
(define (generate-adt names)
  (lambda (arg1 . args)
    (lambda (message)
      (list-ref (cons arg1 . args) (position message names))))))
```

And another, more lengthy solutions:

```
(define (generate-adt names)
  (define (class-like . args)
    (define (object-like message)
      (list-ref args (position message names)
        object-like)
      class-like)
  class-like)
```

Solutions that didn't use LIST-REF and POSITION but still worked also received credit.

Grading: 8 points total.

- \* 4 for recognizing and formalising the domain and range of

- generate-adt correctly, i.e. writing (lambda args (lambda (message) ...))

- 2 points for writing two lambdas matching the arguments in the right order

- 2 points for writing args without parenthesis or writing (arg1 . args)

- (and failing to use it correctly)

- \* 4 points for the body (list-ref args (position message names))

- 2 points for composing list-ref and position

- 2 points for using the right arguments to list-ref and position

- 0 points subtracted for swapping the list and the index argument to list-ref

Question 11:

Write a function PATH-ACCUMULATE that, given a Tree, a function of two arguments, and a base case for that function, returns a new Tree in which the datum of every node is the accumulation of the path



in the original Tree from  
the root to that node.

```
(define (path-accumulate fn base tree)
  (let ((new-datum (fn (datum tree) base)))
    (make-tree new-datum
                (map (lambda (child) (path-accumulate
                                     fn new-datum child))
                     (children tree) ))))
```

This was a pretty straightforward tree recursion question. As with most such problems, it could be done either using MAP, or a separate

FOREST-ACCUMULATE helper function. The basic idea of such a question is

"we're making a new Tree, so what's the new datum, and what should I do to the children?"

In this case, the new datum is what you get by combining the old datum

with the running accumulation -- which is BASE. The children aren't

filtered in any way, just PATH-ACCUMULATED versions of the original

children. However, you have to remember to use the new datum as the new

base as well -- otherwise there's not actually any accumulation!

Scoring: 8 points total

\* 8 - perfect

\* 7 - trivial errors, such as using + and 0 instead of FN and BASE

\* 6 - correct, but forgot to use the new base in the recursive call

\* 5 - a good attempt at Tree recursion, plus applying FN to the datum and BASE

\* 2 - applying FN to the datum and BASE, but no clear tree recursion

\* 0 - other

## Question 12:

There were four key parts to this question. The first was realizing that the input to LOOKUPSTATIC was going to be a procedure /name/, not an actual procedure. The second was knowing that the procedure statics were stored in a frame inside the procedure (once you had looked it up), as specified in project 4. The third was knowing how to get the value out of the statics frame. Last was knowing how to add LOOKUPSTATIC as a Logo primitive.

One correct implementation:

```
(define (lookupstatic proc-name var-name)
  (let ((proc (lookup-procedure proc)))
    (lookup-variable-value var-name (list (statics-frame proc)))))
```

The calls to LOOKUP-PROCEDURE and STATICS-FRAME were pretty much the same for everybody. The trouble came with getting the variable value /out/ of the statics frame. Many people forgot that LOOKUP-VARIABLE-VALUE takes an /environment/, not a frame.

There were three ways around this that came up in answers:

1. Like the sample solution above, make an environment containing a single

frame. Probably an even better solution would have been to use

```
(adjoin-frame (statics-frame proc) the-empty-environment),
```

and a few people did do this. But we have said an environment is just a

list of frames, not a proper ADT.

2. Copy the SCAN procedure from LOOKUP-VARIABLE-VALUE, and tweak it for use here.

3. Use the FRAME-VARIABLES and FRAME-VALUES selectors to extract the variable names and values, then use POSITION and LIST-REF as in Question 10.

Last, you had to add your procedure to Logo:

```
(add-prim 'lookupstatic 2 lookupstatic)
```

Grading: 8 points total

- \* 1 for the call to ADD-PRIM, all-or-nothing
  - No points off for using a count of (2), as long as your procedure actually takes an environment as a parameter.
- \* 7 for the implementation of LOOKUPSTATIC
  - 7 - perfect
  - 5 - treated (statics-frame proc) as an environment instead of a frame
  - 4 - attempted to look up the procedure, extract the statics frame from the procedure, and look up the variable in the statics frame.
  - 2 - only did one or two of the above.
  - 0 - anything else