

CS 61B: Data Structures (Autumn 2010)

Midterm II

Solutions


Problem 1. (10 points) A Miscellany.

a. Determining whether the graph contains a particular edge (x, y) . Other acceptable answers include inserting an edge and deleting an edge. A particularly interesting solution, contributed by one student on his exam, is finding all the edges directed *into* a particular vertex. This takes $O(|V|)$ time with an adjacency matrix, and $O(|V| + |E|)$ time with an adjacency list.

b.

✖	3	5	4	7	6	5	7	9	8	
✖	1	3	4	7	5	5	7	9	8	6
✖	3	5	4	7	6	5	7	9	8	

c.



d. cdadcda.

e. Three acceptable answers (each worth full points):

If any one of the characters is zero, the hash code is zero.

Anagrams all have the same hash code.

A string's hash code doesn't change if you insert characters equal to one.

Problem 2. (8 points) Algorithms.

a, b. With `bottomUpHeap()`, create a heap containing the n numbers. This takes $\Theta(n)$ time. Then, call `removeMin()` j times, taking $\Theta(j \log n)$ worst-case time. The total worst-case running time is in $\Theta(n + j \log n)$.

An alternative solution uses a *max-heap*, which reverses the heap-order property, so that the greatest key is at the top of the heap. Starting from an empty max-heap, insert j keys (it doesn't matter which ones) into the heap. Then repeat the following operations for each remaining key: insert the key into the heap, and remove the maximum key from the heap. After you've processed all the keys, the heap contains the j smallest. At worst, this takes $\Theta(n \log j)$ time.

The optimal solution—which requires knowledge you won't learn until later in the semester—is to use a median-finding algorithm to find the j th-smallest key in $\Theta(n)$ expected time. Then, a linear scan through the array can find the other $j - 1$ smallest items, also in $\Theta(n)$ expected time.

c. Output-sensitive.

d. Starting from the root, walk down the tree. When you encounter a key less than k , go right. When you encounter a key greater than or equal to k , go left. Repeat until you reach a null pointer. (A useful intuition is that this is equivalent to performing a `find` operation on the key $k - \epsilon$, where ϵ is infinitesimal.)

As you go down the tree, keep track of the largest key you've encountered so far that's less than k . When you reach a null pointer, return that value. The tree might not contain any key less than k , in which case you can return `null` or some other value that means “no such key.”

Problem 3. (7 points) **Sibling-based tree height.**

The main idea is that if `this` node has a next sibling, then `height()` should return the *maximum* height of `this` node and its siblings to the right. This information allows the parent to calculate its height. The other idea is to use a postorder traversal to calculate the heights.

```
public class SibTreeNode {
    public Object item;
    public SibTreeNode parent;
    public SibTreeNode firstChild;
    public SibTreeNode nextSibling;

    int height() {
        int h = 0;
        if (firstChild != null) {
            h = 1 + firstChild.height();
        }
        if (nextSibling != null) {
            int hSib = nextSibling.height();
            if (hSib > h) {
                return hSib;
            }
        }
        return h;
    }
}
```